

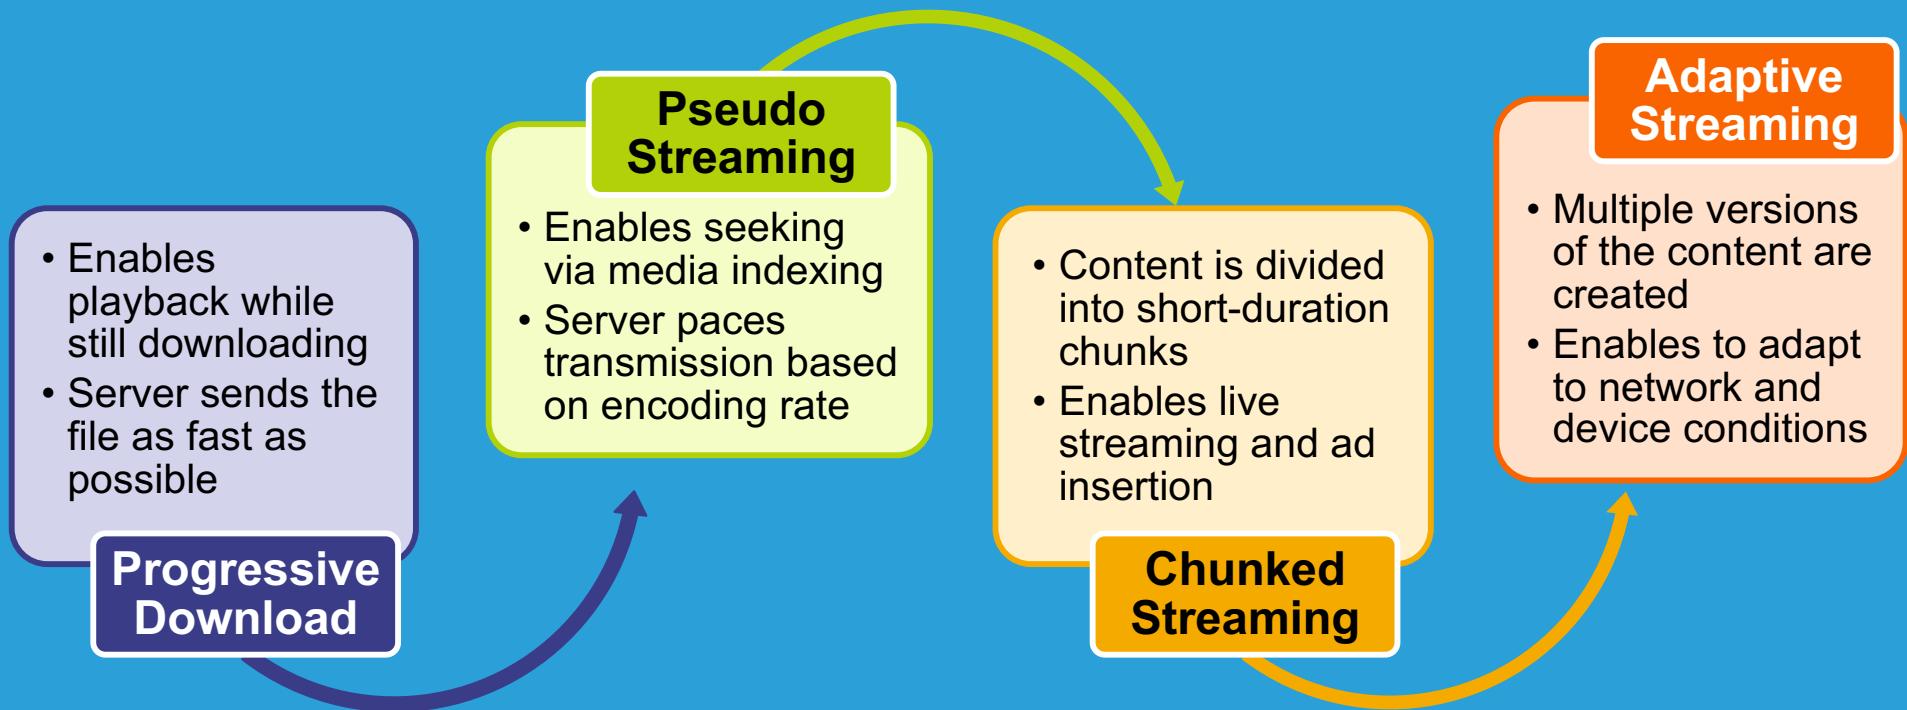
# ANATOMY OF A DASH CLIENT

- 1. TASKS**
- 2. SCALING**
- 3. ANALYTICS**
- 4. COOL FEATURES**

ALI C. BEGEN, PH.D.  
[HTTP://ALI.BEGEN.NET](http://ali.begen.net)



# Video Delivery over HTTP

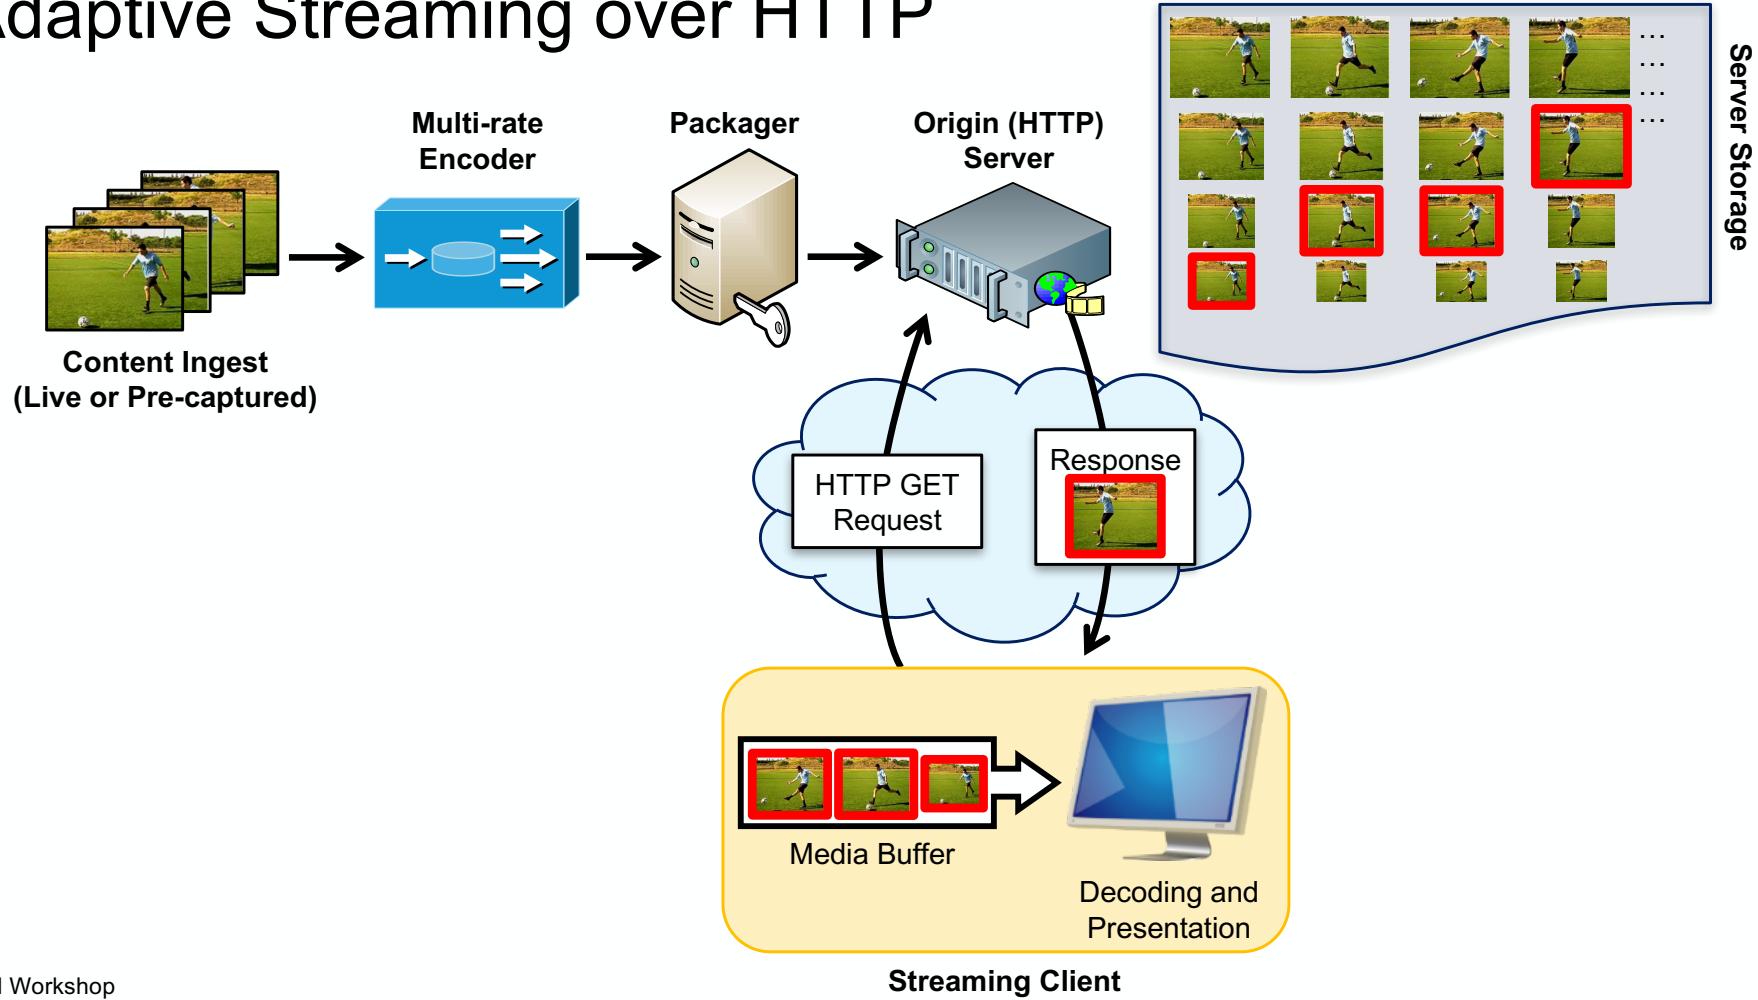


# Dead, Surviving, Maturing and Newborn Technologies

- **Move Adaptive Stream (Long gone)**
  - <http://www.movenetworks.com>
- **Microsoft Smooth Streaming (Legacy)**
  - <http://www.iis.net/expand/SmoothStreaming>
- **Adobe Flash (Almost dead)**
  - <http://www.adobe.com/products/flashplayer.html>
- **Adobe HTTP Dynamic Streaming (Legacy)**
  - <http://www.adobe.com/products/httpdynamicstreaming>
- **Apple HTTP Live Streaming (The elephant in the room)**
  - <http://tools.ietf.org/html/draft-pantos-http-live-streaming> (Soon to be an RFC)
- **MPEG DASH and CMAF (The standards)**
  - <http://mpeg.chiariglione.org/standards/mpeg-dash>
  - <http://mpeg.chiariglione.org/standards/mpeg-a/common-media-application-format>

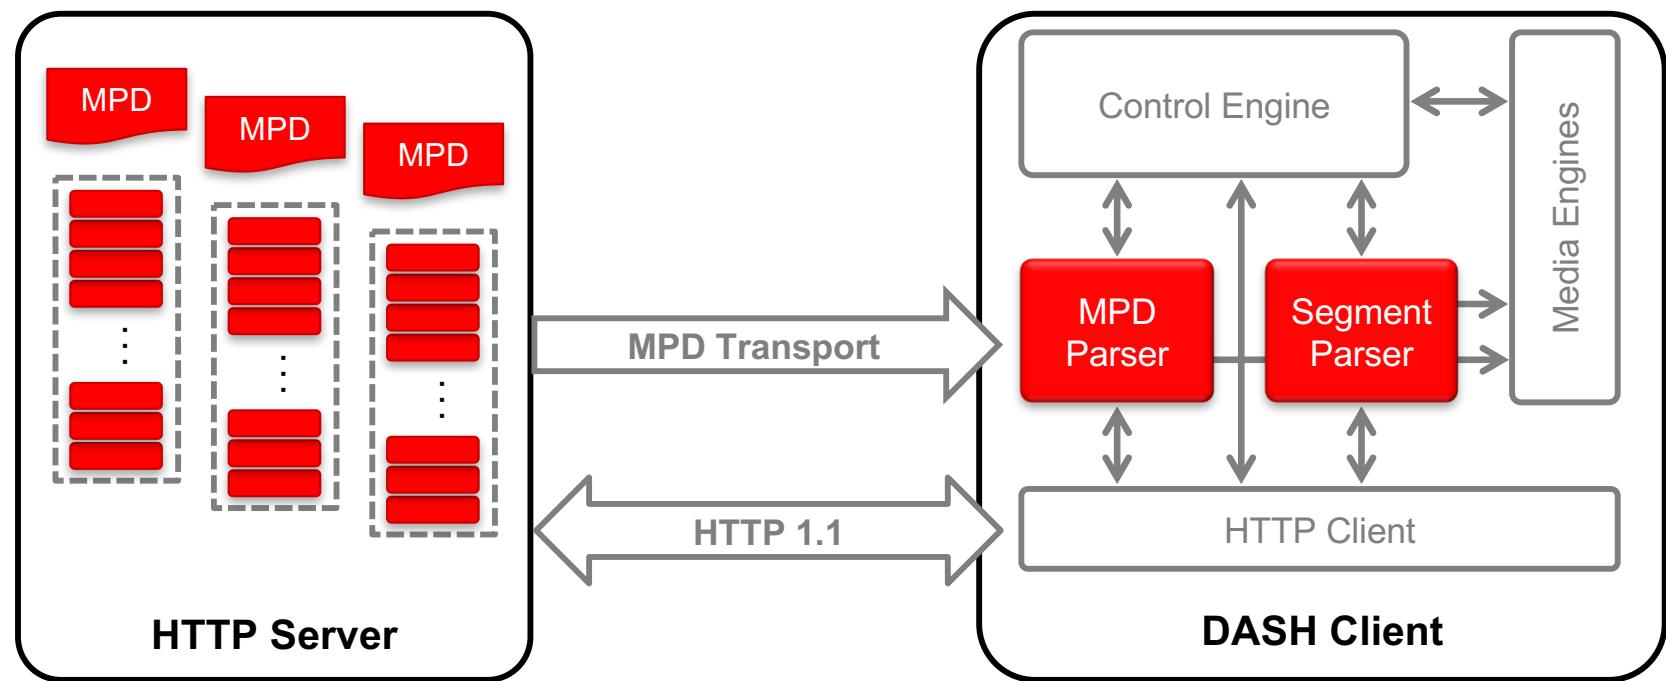


# Adaptive Streaming over HTTP



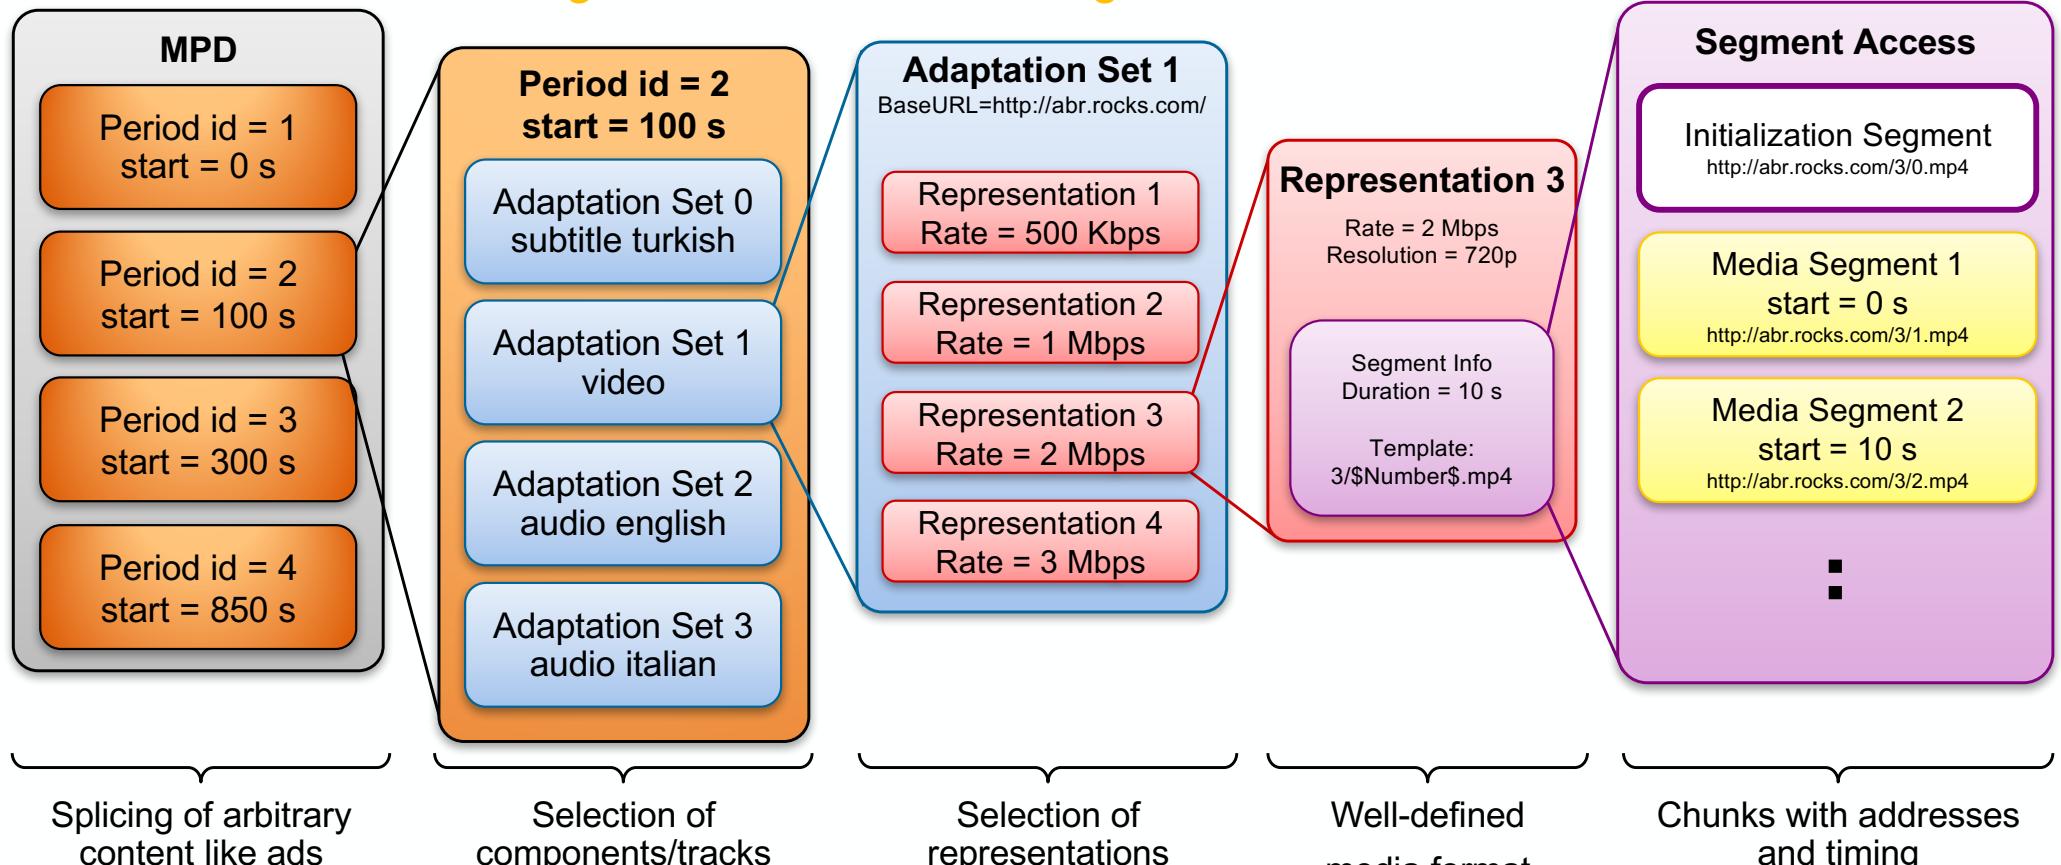
# Scope of MPEG DASH

Shown in Red



# DASH Manifest – Media Presentation Description (MPD)

## List of Accessible Segments and Their Timings



# Smart Clients

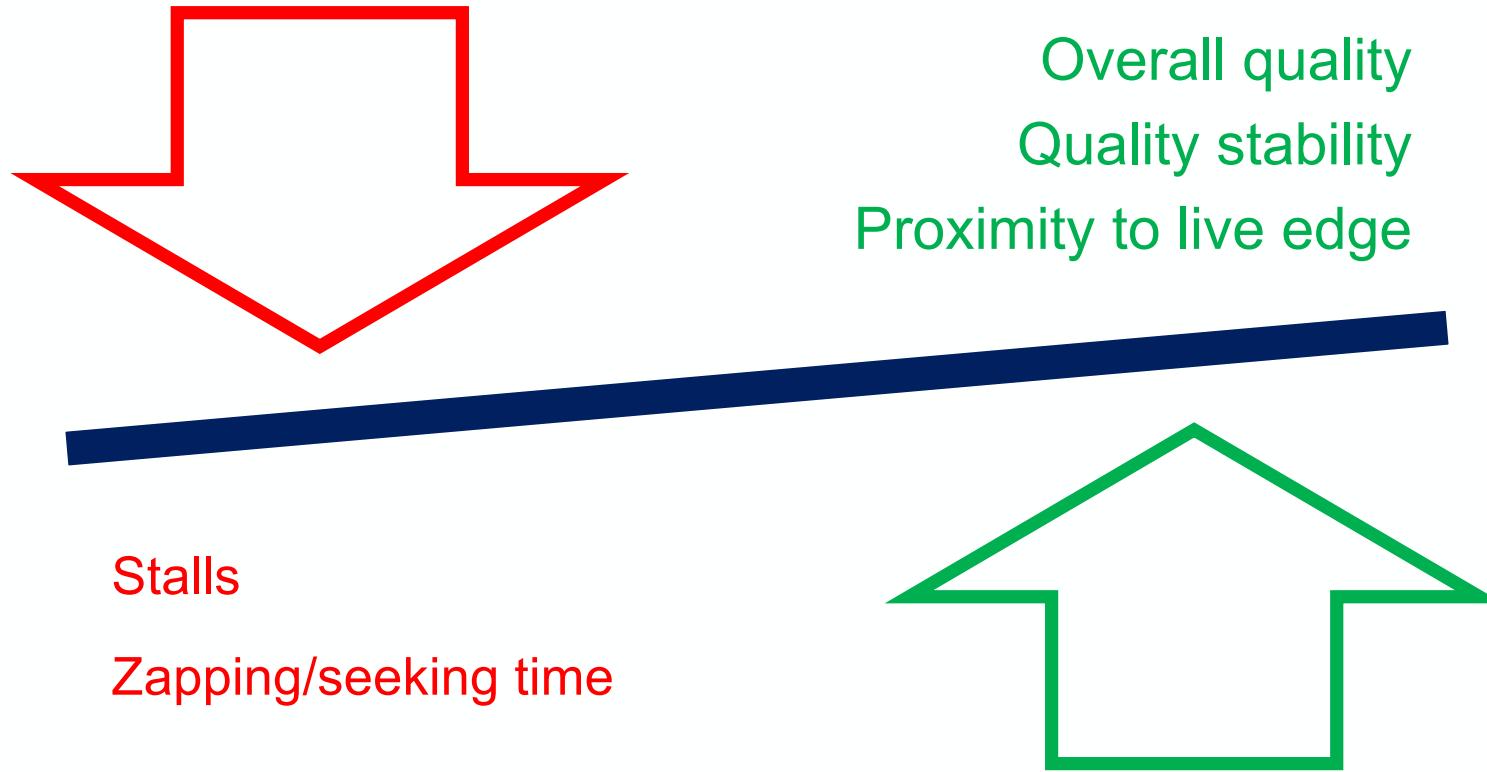
HTTP Server



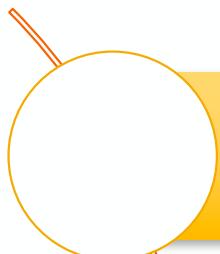
Client

- Client fetches and parses the manifest
  - Client uses the OS-provided HTTP stack (HTTP may run over TCP or QUIC)
  - Client uses the required decryption tools for the protected content
- 
- Client monitors and measures
    - Size of the playout buffer (both in bytes and seconds)
    - Chunk download times and throughput
    - Local resources (CPU, memory, window size, etc.)
    - Dropped frames
- Client performs adaptation**
- Client measures and reports metrics for analytics

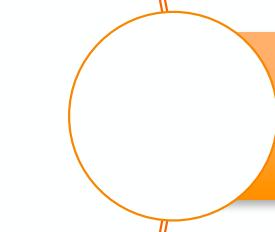
## Tradeoffs in Adaptive Streaming



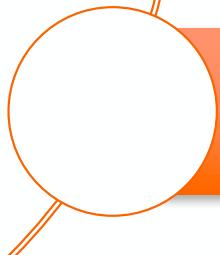
# Types of Browser-Based Playback



Type 1: Minimum control architectural model for HTML5 support of adaptive streaming where manifest and heuristics are managed by the user agent



Type 2: Adaptation control architectural model for HTML5 support of adaptive streaming providing script manageable features

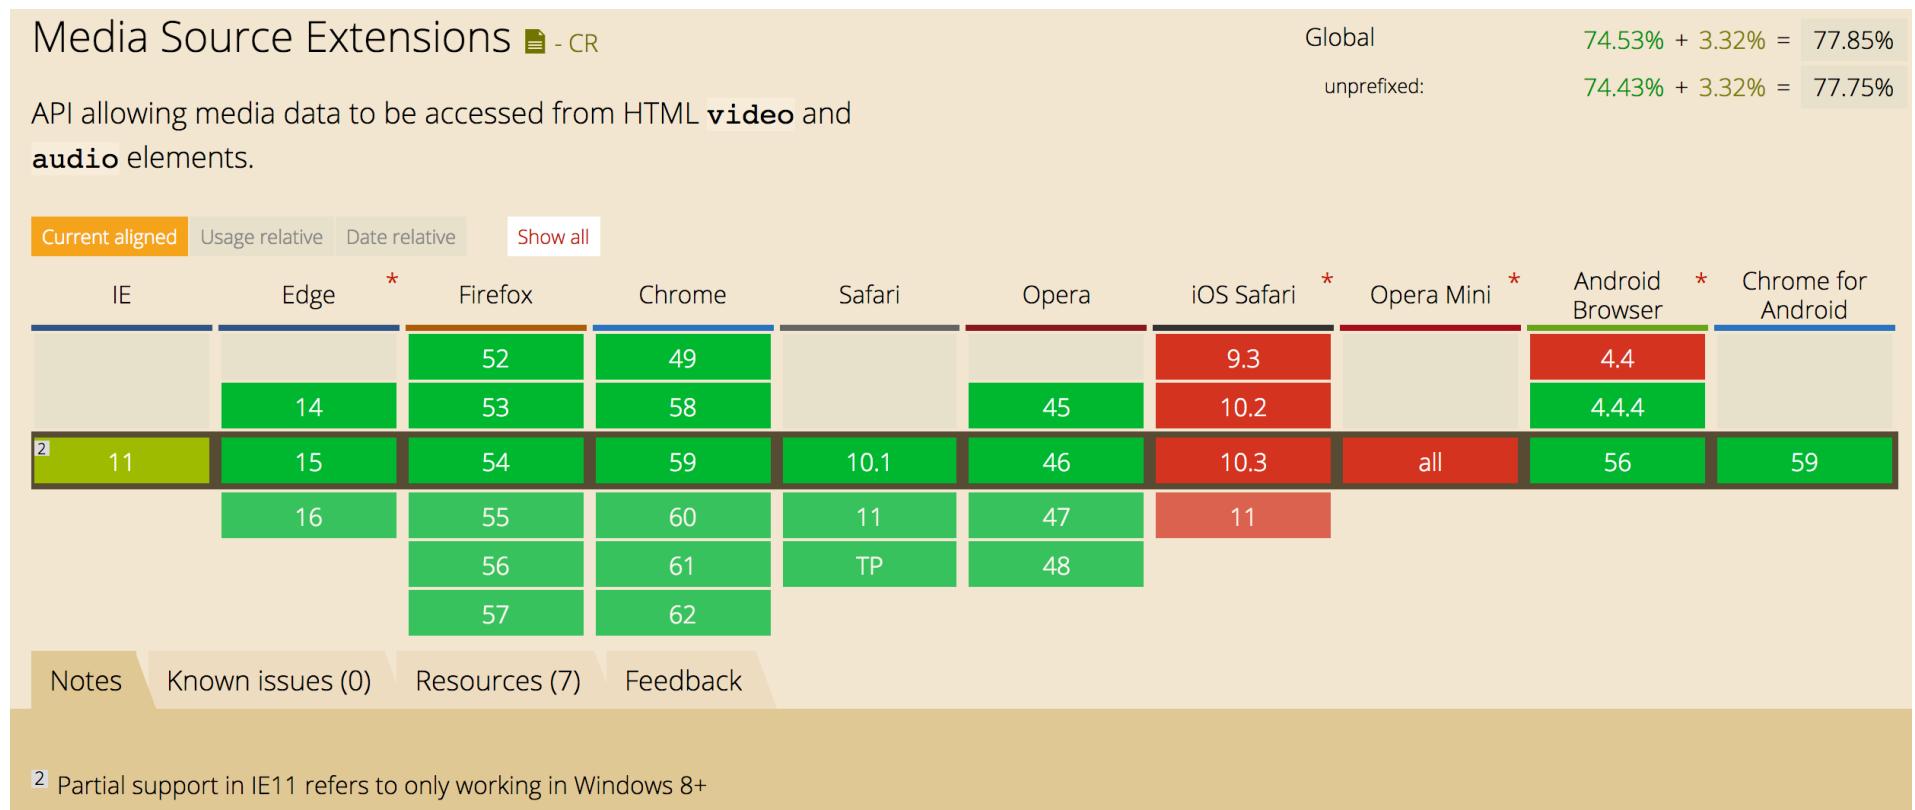


Type 3: Full media control architectural model for HTML5 support of adaptive streaming allowing script to explicitly send the media segments (MSE + EME)

Source: CTA WAVE

DASH Workshop

# MSE Support in Web Browsers



Source: <http://caniuse.com/#search=mse>

# DRM Support on Desktop Browsers

Browser	OS	EME/CDM	Flash Player	NPAPI/ Silverlight 5
Chrome	Win	Widevine	(Yes)	No
	OS X		(Yes)	No
	Linux		(Yes)	No
Firefox	Win	Widevine	Yes	No
	OS X		Yes	No
	Linux		Yes	No
Safari	> OS X Yosemite	(FairPlay)	Yes	Yes
	< OS X Yosemite	No	Yes	Yes
IE/Edge	< Win 7	No	Yes	Yes
	Win 10	PlayReady	Yes	No

# DRM Support on Mobile Platforms

Platform	Browser	Native App / DRM
iOS	No MSE/EME yet HLS (AES-128 CBC) via <video>	Native SDK and WebView: FairPlay
Android	MSE/EME	Native + Android SDK (MediaDRM APIs) (Widevine + OMA v2) ExoPlayer
Windows 10	MSE/EME	Native + WebView with Widevine WebViews/Hosted Web Apps with PlayReady

## Streaming over HTTP – The Promise

- Leverage tried-and-true Web infrastructure for scaling
  - Video is just ordinary Web content!
- Leverage tried-and-true TCP
  - Congestion avoidance
  - Reliability
  - No special QoS for video

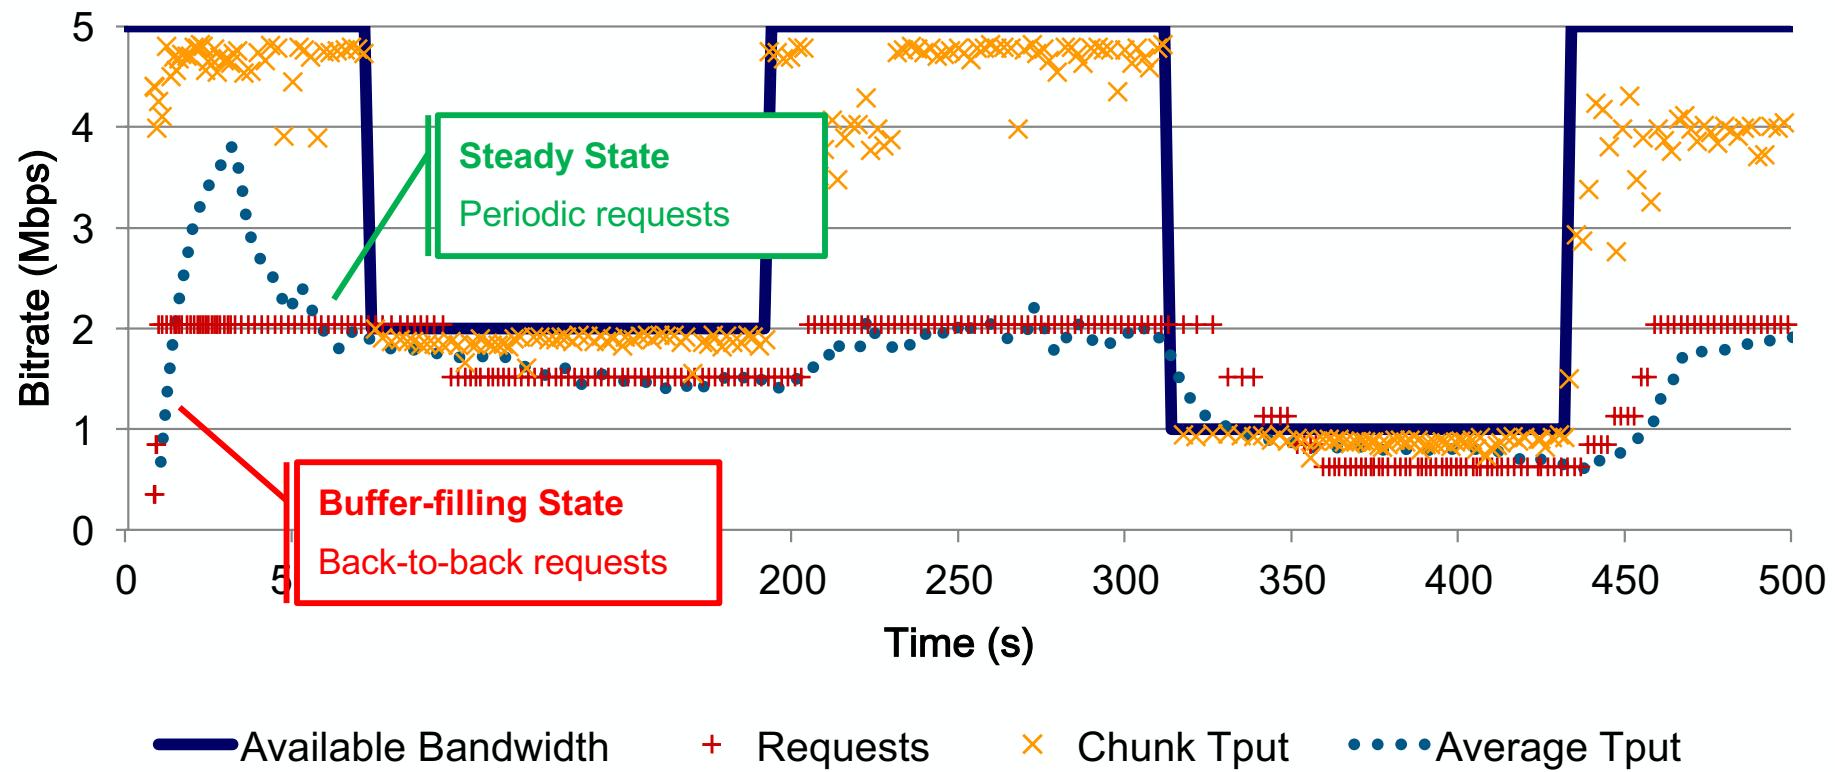
*It should all “just work”* ☺

## Does It Just Work?

- When a streaming client competes with other types of traffic, mostly yes
- When streaming clients compete with each other, we begin to see problems:
  - The clients' adaptation behaviors interact with each other
  - The competing clients form an “accidental” distributed control-feedback system
- Unexpected behaviors will result in places like
  - Multiple screens within a household
  - ISP access and aggregation links
  - Small cells in stadiums and malls

# Demystifying a Streaming Client

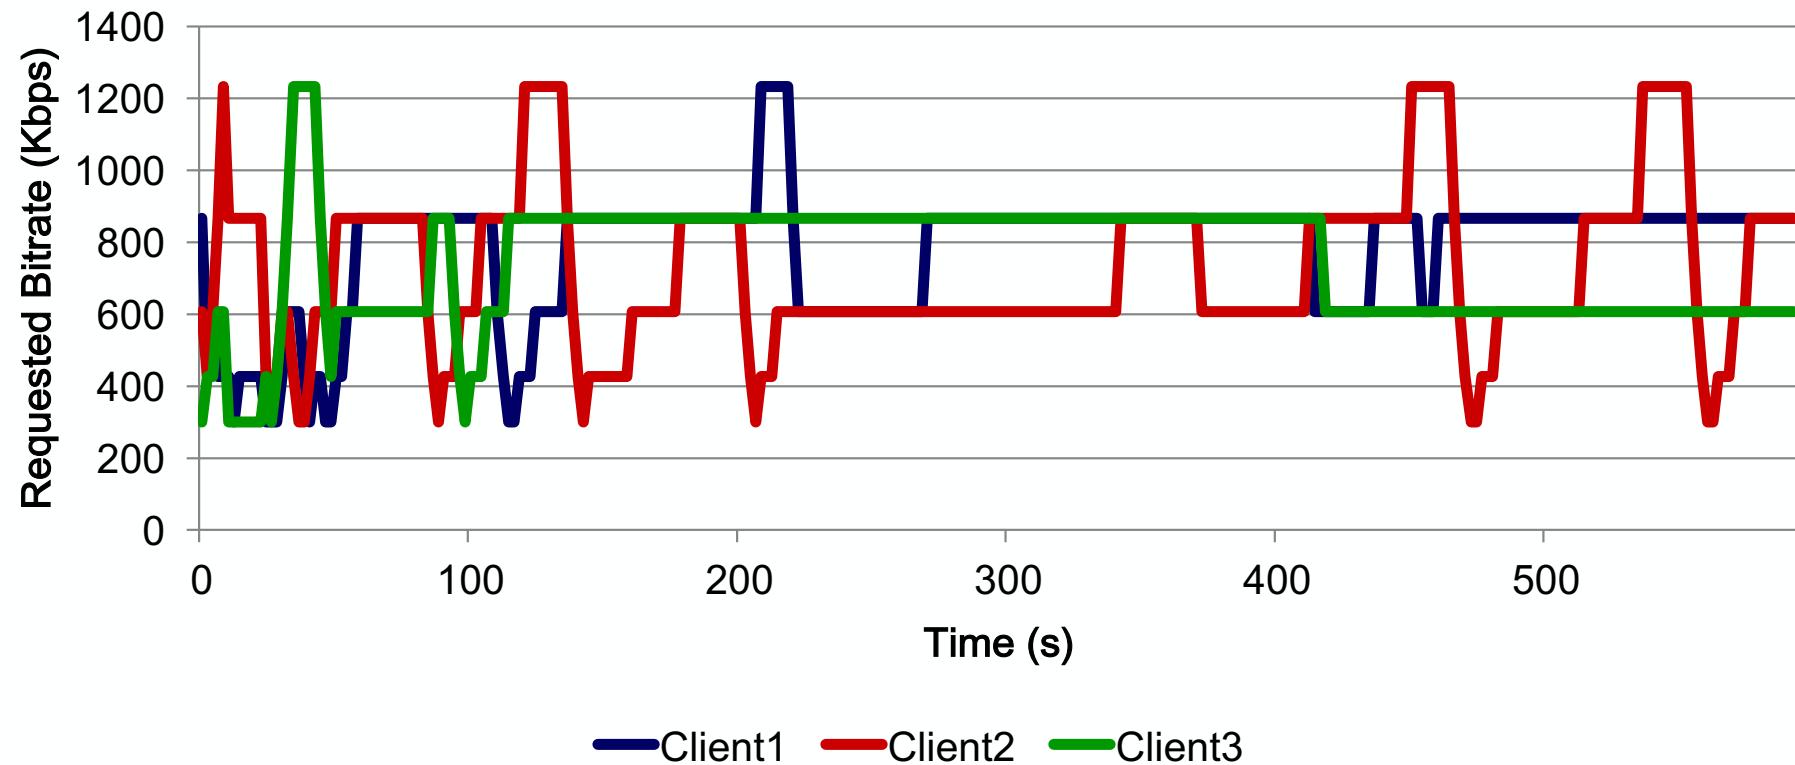
A Single Microsoft Smooth Streaming Client under a Controlled Environment



Reading: "An experimental evaluation of rate-adaptation algorithms in adaptive streaming over HTTP," ACM MMSys 2011

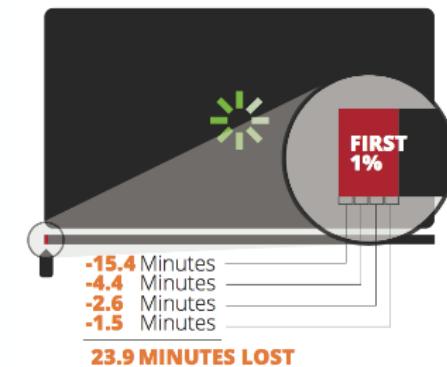
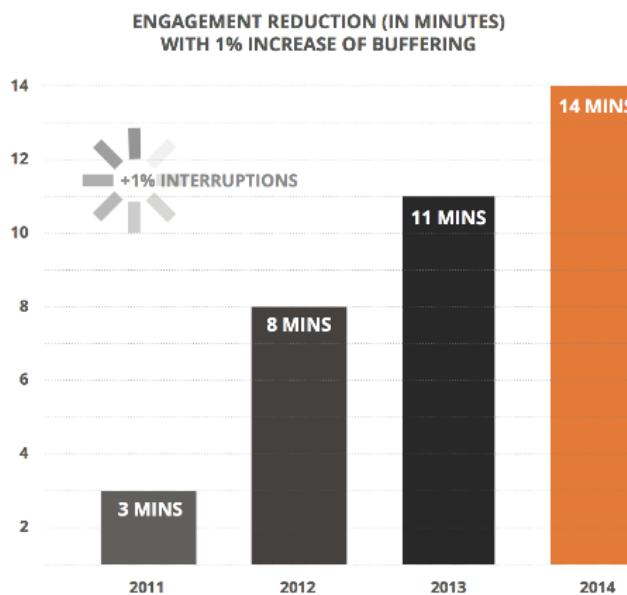
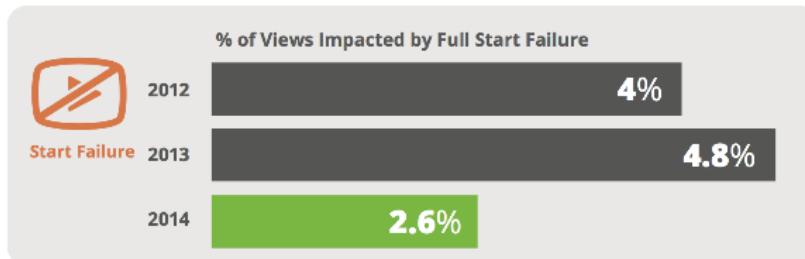
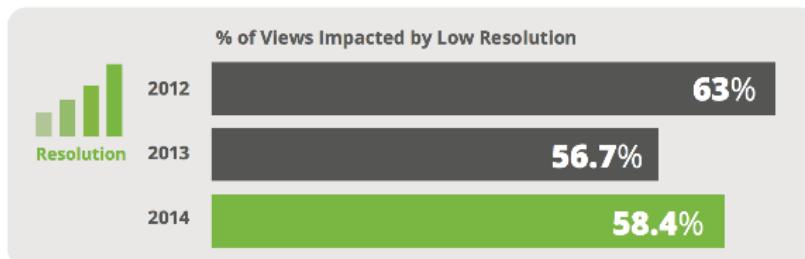
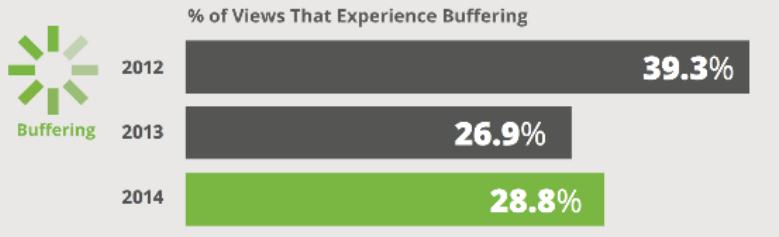
# A Simple Test Scenario

10 (Commercial) Streaming Clients Sharing a 10 Mbps Link



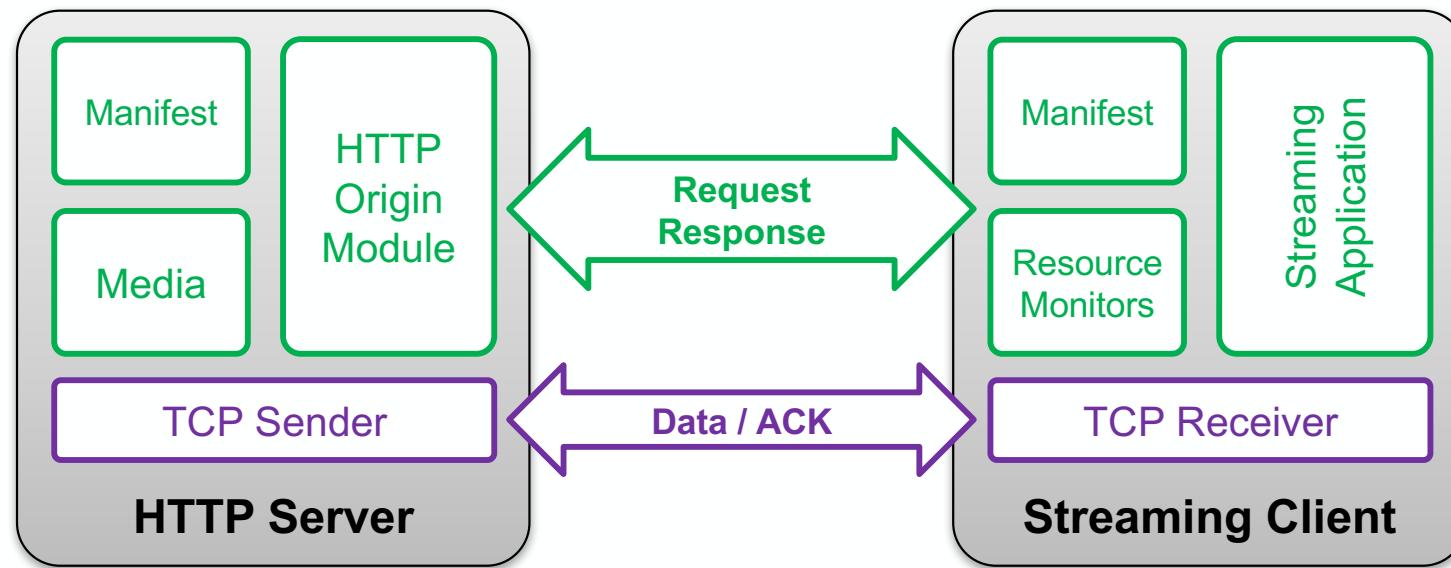
# Selfishness Hurts Everyone

## Viewer Experience Statistics



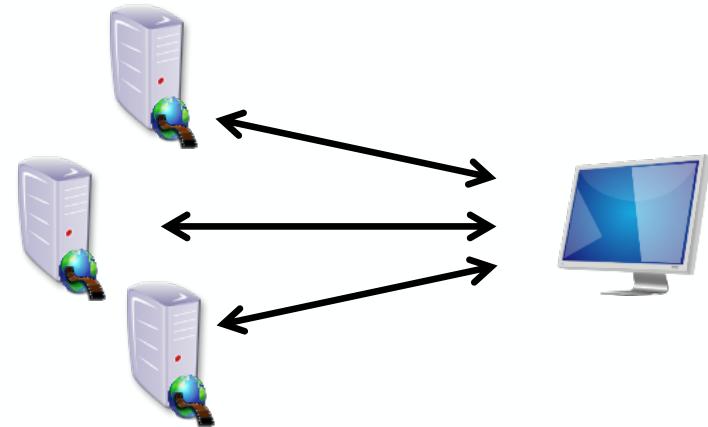
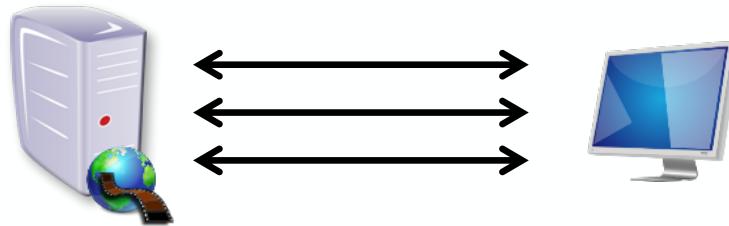
Source: Conviva Viewer Experience Report, 2015

# Inner and Outer Control Loops



There could be multiple TCPs destined to the same or different servers

# Streaming with Multiple TCP Connections



- Using multiple concurrent TCPs
  - Should not be used to greedily get a larger share of the bandwidth
  - Can help mitigate head-of-line blocking
  - Allows fetching multiple (sub)segments in parallel
  - Allows to quickly abandon a non-working connection without having to slow-start a new one

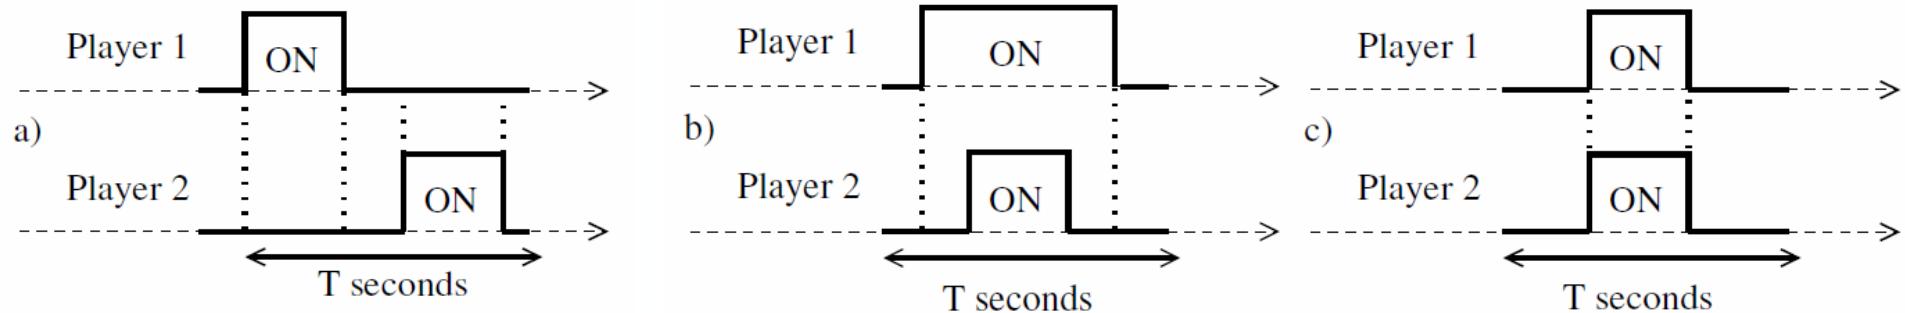
System performance deteriorates very quickly if many clients adopt this approach without limiting the aggregated bandwidth consumption

# Understanding the Root Cause

## Two Competing Clients

- Depending on the timing of the ON periods:
  - Unfairness, underutilization and/or instability may occur
  - Clients may grossly overestimate their fair share of the available bandwidth

*Clients cannot figure out how much bandwidth to use until they use too much*



Reading: "What happens when HTTP adaptive streaming players compete for bandwidth?," ACM NOSSDAV 2012

# How to Solve the Issues?

Fix the clients  
(Use a better algorithm like PANDA or BOLA)

## Solution Approaches

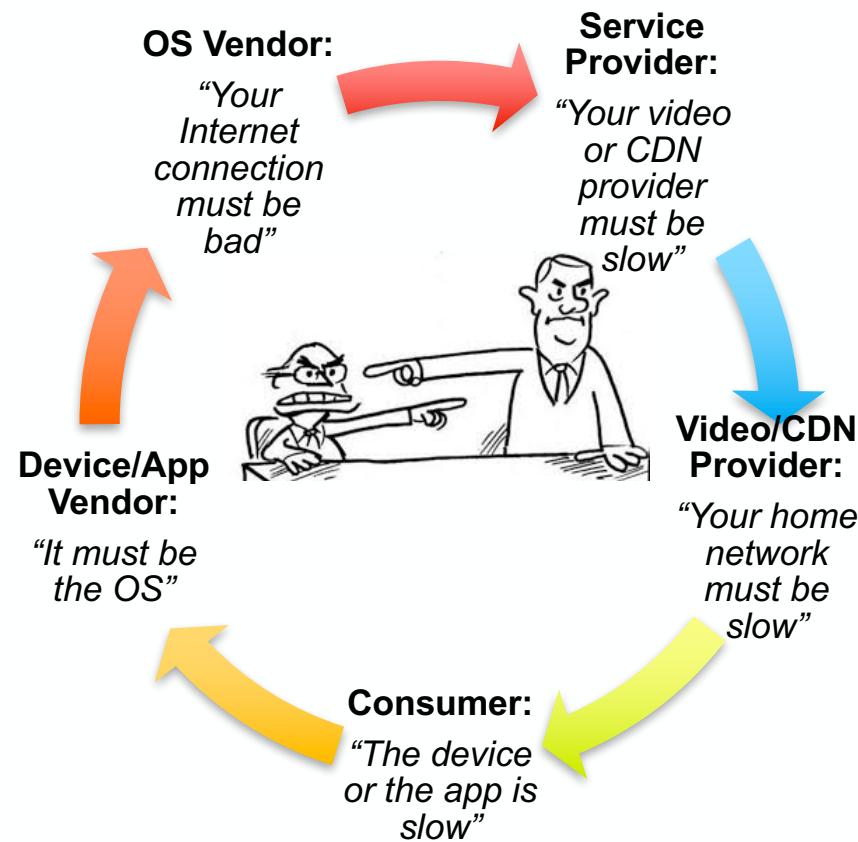
Get support from the network  
(QoS in the core/edge, SDN, etc.)

Enable a control plane  
(E.g., 23009-5 SAND)

# Commercial OTT Issues

- Content format issues
  - Each asset is copied to multiple media formats
    - different video codecs
    - different audio codecs
    - different (regional) frame rates
  - Cost to content creators and distributors
  - Inefficiencies in CDNs and higher storage costs
- Platform issues
  - Lack of consistent app behavior across platforms
  - Varying video features, APIs and semantics across platforms
- **Playback issues**
  - **Codec incompatibility**
  - **Partial profile support**
  - **Switching bitrate glitches**
  - **Audio discontinuities**
  - **Ad splicing problems**
  - **Long-term playback instability**
  - **Request protocol deficiencies**
  - **Memory problems, CPU weaknesses**
  - **Scaling (display) issues**
  - **Variable HDR support**
  - **Unknown capabilities**

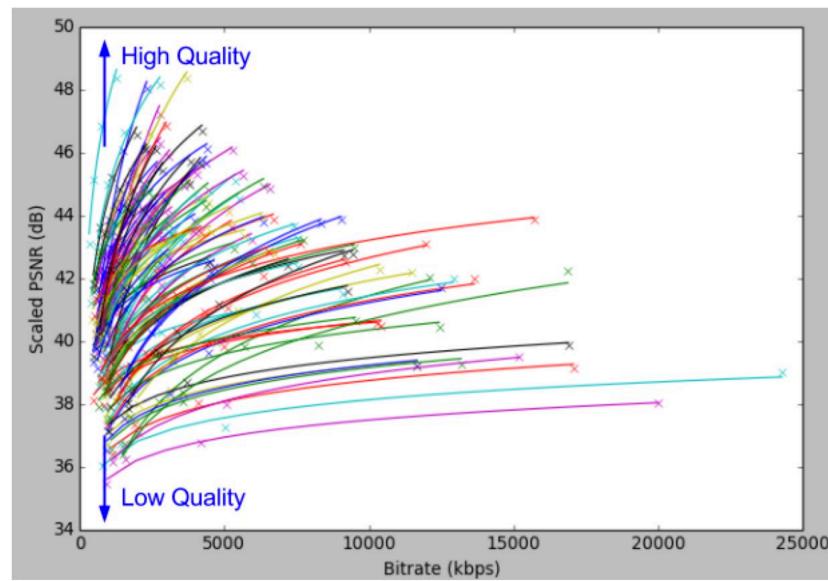
# That is Why We Need Analytics



# Netflix's Content-Based Encoding Method

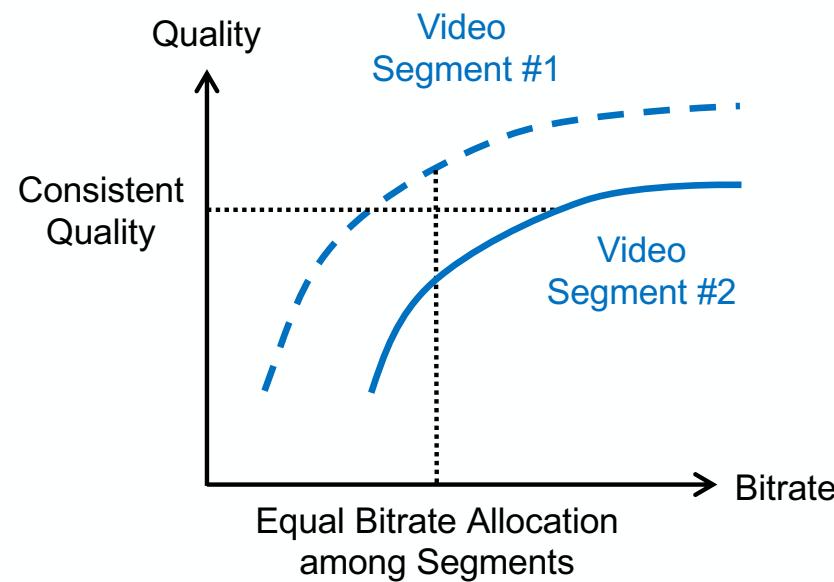
**Increased bandwidth brings new opportunities**

Drives overall maturity of IP/ABR delivery technologies, introduces new opportunities –  
Per Title Encode Optimization



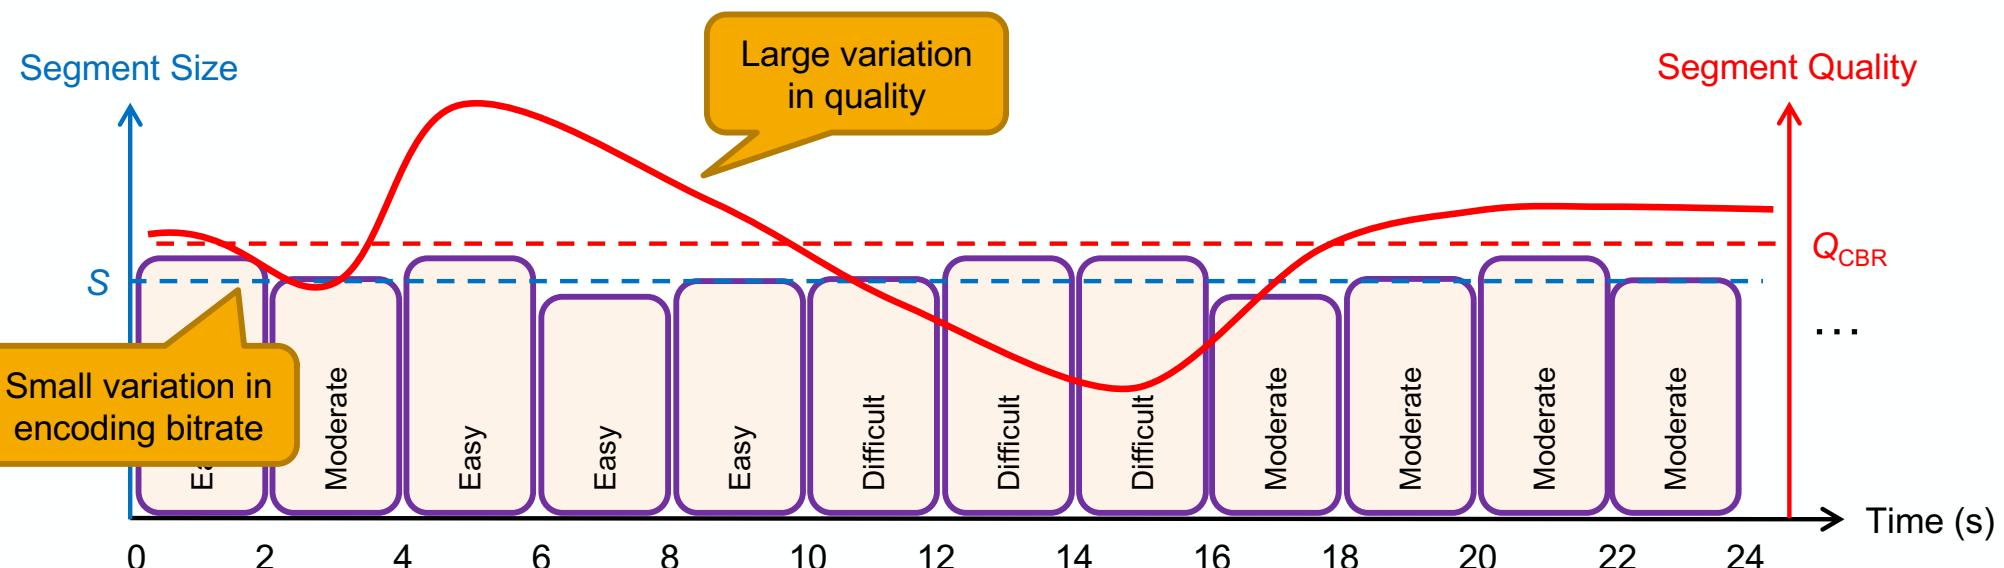
Source: Netflix Tech Blog, December 2015

# Segments Have Different Complexities



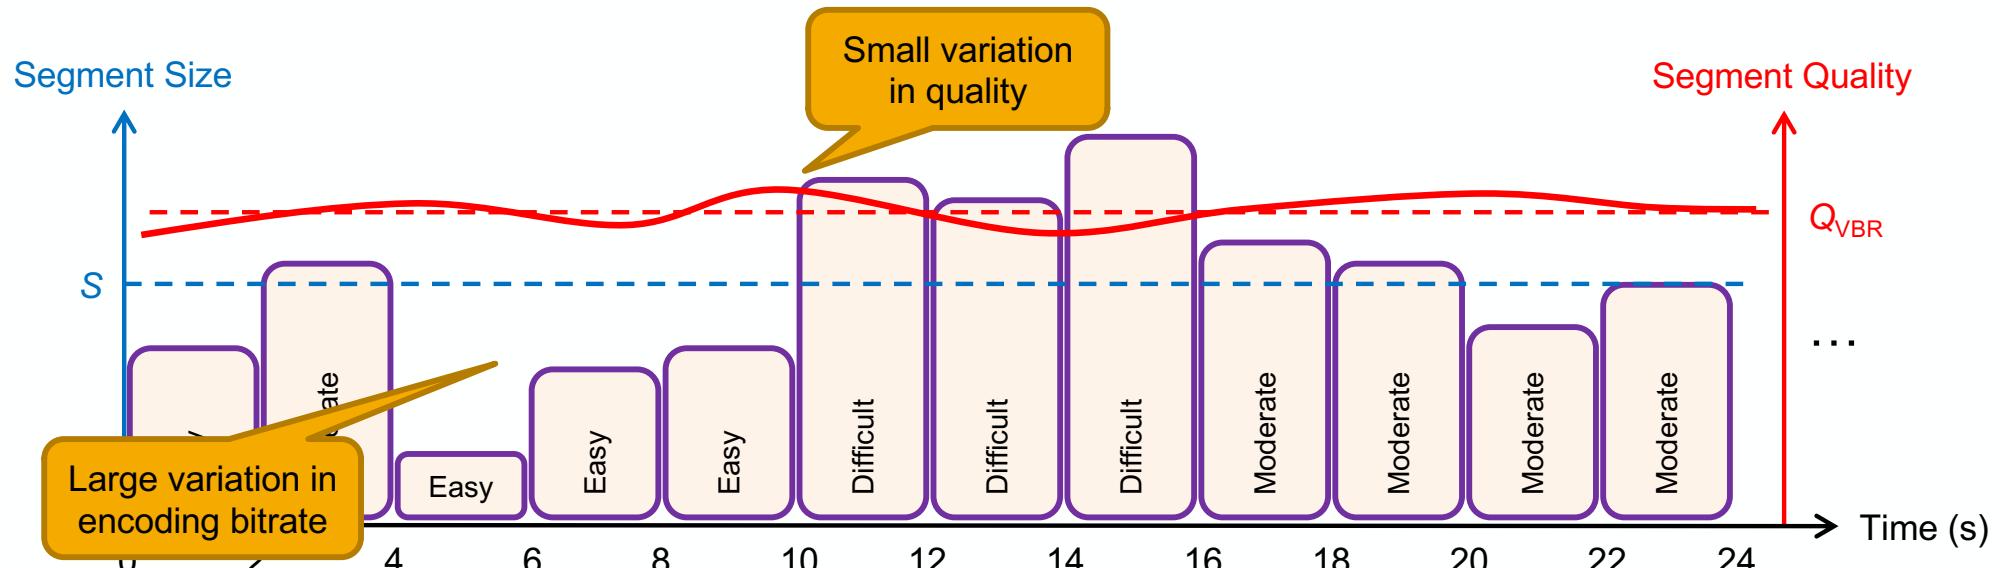
# Adaptation Feature Does Not Deliver Consistent Quality

Guidelines Limited Bitrate Variability to (Mostly) 10% So Far



If there is something worse than having to watch a video at a lousy quality, it is to watch that video with varying quality

# What if We Encode in a More Subtle Fashion?

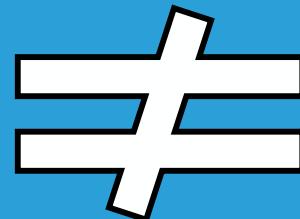


While we spend the same total amount of bits, we not only increase average quality but also reduce quality variation

Note: HLS authoring spec for ATV allows 2x capping rate for VoD. For linear content, variability is limited to 10-25% range.

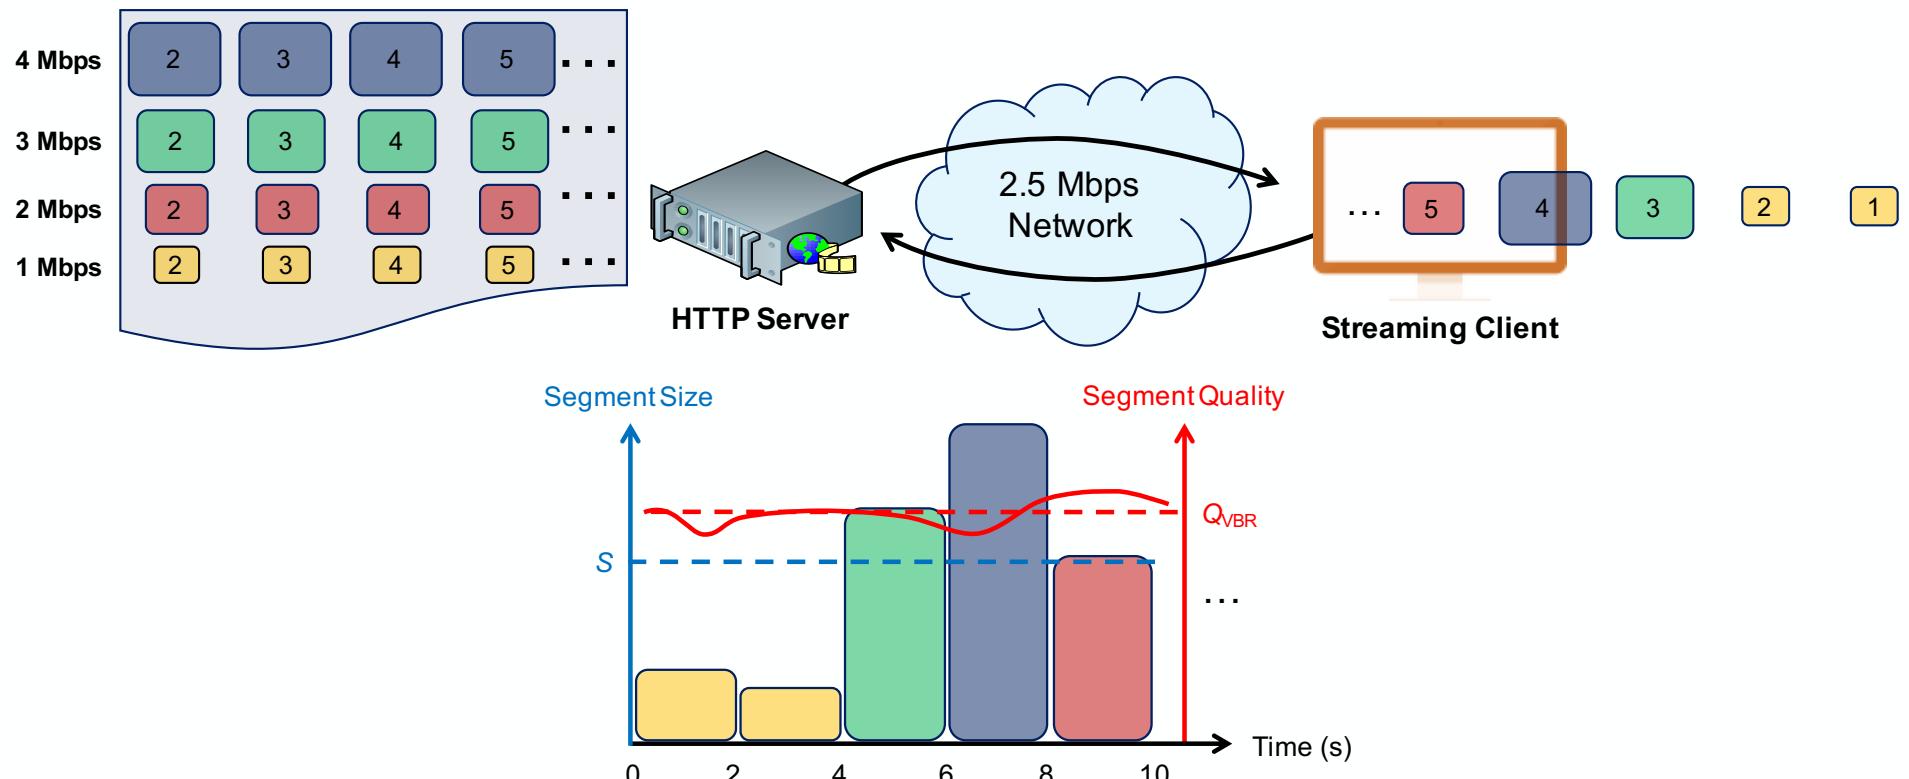
**Generating VBR-encoded segments is easy,  
but streaming them is not!**

Content Aware  
Encoding



Content Aware  
Streaming

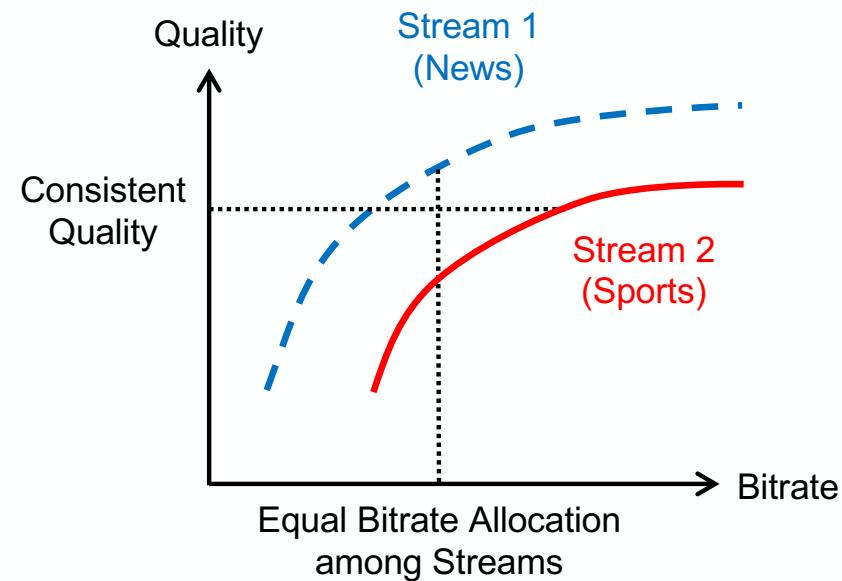
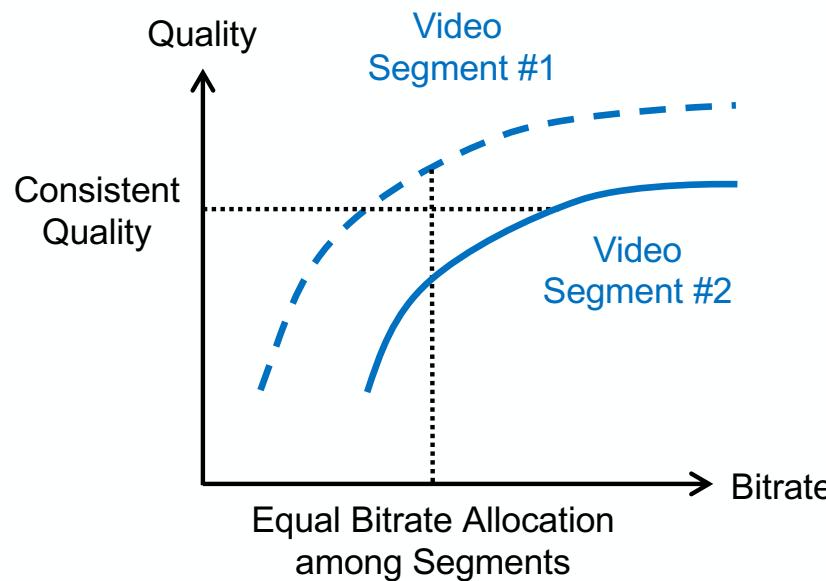
# Multiple Representations Naturally Enable “Cherry-Picking”



Reading: “Streaming video over HTTP with consistent quality,” ACM MMSys 2014

# Dimensions – In-Stream vs. Across-Streams

- Same principle applies to both:
  - In-stream: Temporal bit shifting between segments
  - Across-streams: Bit shifting between streams sharing a bottleneck link



Reading: “Spending quality time with the Web video,” IEEE Internet Comput., 2016



**ANY QUESTIONS?**